

Christopher Anaya

Huntsville, AL | Open to remote or relocation
chris.anaya.ai@gmail.com | +1.938.257.5842 | chranama.github.io

SUMMARY

Full-stack AI product engineer building React, Go, Python, and Postgres systems that turn model output into reviewable product workflows. Built MealCheck, a hosted verifier with local-model normalization, source-linked review, deterministic checks, evaluation gates, and artifact-backed reports.

EXPERIENCE

Open-Source Software Engineer

Mar 2026 - Present

- Built MealCheck as a deployed React/Vite product with a Go API/worker, Postgres run state, filesystem artifacts, Cloudflare frontend/tunnel deployment, and private llama.cpp normalization.
- Shipped source-linked review, correction, rejection, rewrite, progress, deletion, report, recommendation, and source-inspection flows so users can inspect and recover from model normalization before deterministic checks run.
- Added normalization and food-resolution evaluation gates, live local-model smoke paths, model-comparison artifacts, Python operator summaries, and Go/TypeScript CI checks to make release behavior reproducible.

Graduate Student Researcher

Oct 2023 - Jun 2025

LASIGE / University of Lisbon

- Built a fixed-seed biomedical QA fine-tuning and evaluation workflow with deterministic perturbations, BioASQ-style scoring, and saved clean-vs-perturbed artifacts for model robustness analysis.
- Packaged the thesis workflow as BioLLM-Finetune with manifests, tables, figures, validation metadata, and reproducible report artifacts; work supported a CLEF 2024 workshop publication.

SELECTED PROJECTS

MealCheck

- Deployed a hosted verifier that accepts one concise ingredient-level day of meals, normalizes with a server-owned local model, pauses for review/correction, then runs deterministic food, nutrient, allergen, and guideline checks.
- Maintains reports with source rows, normalized rows, repairs, unresolved-food summaries, recommendations, and guideline citations; current strict food-resolution gate resolves 885 of 900 FNDDS-grounded items with zero expected-outcome mismatches.

LLM Extraction Platform

- Built schema-constrained sync/async extraction APIs with runtime policy gates, separate worker execution, request-trace artifacts, and Kubernetes deployment artifacts.
- Published eight canonical proof points, CI-backed tests, and evidence manifests that link runtime behavior to inspectable artifacts.

TECHNICAL STRENGTHS

Languages Go, TypeScript/React, Python, SQL/Postgres, HTML, CSS

Software & Tools REST APIs, Vite, JSON Schema, llama.cpp, FastAPI, Docker, Kubernetes, Cloudflare, Redis, Prometheus/Grafana, CI/CD, Git

EDUCATION

- MSc, Data Science | Faculty of Science, University of Lisbon (2022 - 2025)
- MSc, Mathematical Research | University of Valencia (2015 - 2016)
- BA, Astronomy (Minor: Applied Mathematics) | University of Colorado Boulder (2008 - 2013)

PUBLICATIONS

- Anaya, C.; Fernandes, M.; Couto, F.M. *LLM Fine-Tuning With Biomedical Open-Source Data*. CLEF 2024 (BioASQ 12b).